

Demystifying ML, AI & Automation Part I

Spyros Gkezerlis

Oct/Nov 2017

Intro

Spyros Gkezerlis - @sgez

About

BSc Automation Engineering – Telematic Applications

MSc Information Systems – Reinforcement Learning

Interests – Applied ML, RL especially in time domain

Experience

5 years Software Developer in Greece

8 years Engineer and Manager in OTE

3 years Program Manager in DTAG



Ariadni Gkezerli (8 y.o), © 2017

Depiction accuracy: 100%!

Why Lectures?

- Get a better overview of the current landscape in ML, AI & Automation, because they can potentially help us on:
 - Reducing complexity of network
 - Improving experience by Time-to-market, Time-to-repair
 - Repetitive caused costs can be targeted and reduced
 - Forecasting, Automating, Making predictions smarter
- Create a common understanding of what ML, AI & Automation
- “Start with the problem” philosophy
Pros cons, Tools, etc. should not dictate what we should use!
 - Identify what we want to solve
 - Work to the algorithms & models needed
 - Utilise best approach
- Start changing mindset and attitude...

“AI is the New
Electricity”

Andrew Ng

But, be aware of the dark side!

- Keep in mind it can solve particular types of problems!
- ML is not a tool to fix everything
- In some cases needs high degree of mathematics understanding and of course, you should know your problem.



Simplistic Definitions

- Automation Comes from ancient greek word which means the thing that wishes on its own or the has a will or fury by itself
- Artificial Intelligence^{1c} is Human Intelligence Exhibited by Machines
- Machine Learning^{1a} is a field of computer science that gives computers the ability to learn without being explicitly programmed
- Deep learning^{1b} is part of a broader family of machine learning methods based on learning data representations, as opposed to task-specific algorithms.

def·i·ni·tion
defəˈniʃ(ə)n
noun
a statement of the exact meaning of a word,
especially in a dictionary.

1a Machine Learning, Wikipedia Web page, 09 Nov 2017

https://en.wikipedia.org/wiki/Machine_learning

1b Deep Learning, Wikipedia Web page, 09 Nov 2017

https://en.wikipedia.org/wiki/Deep_learning#Definitions

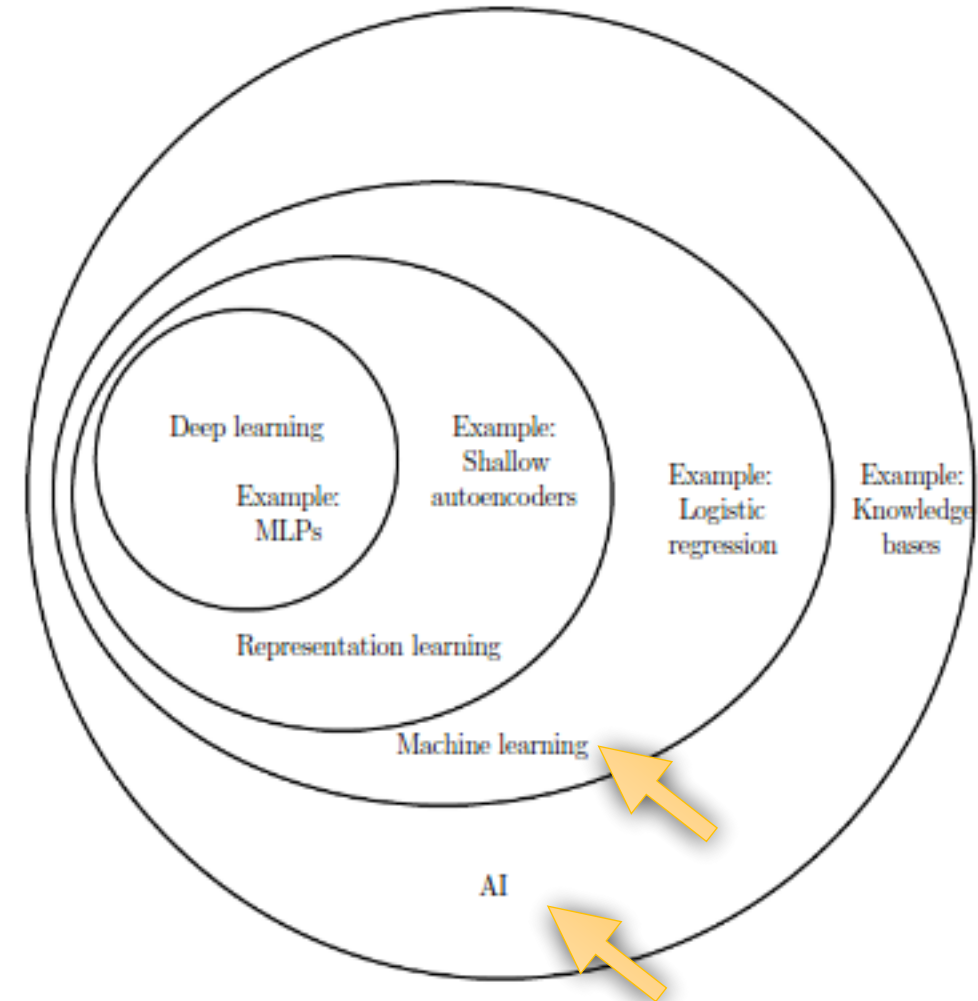
1c What's the Difference Between Artificial Intelligence, Machine Learning, and Deep Learning?, Copeland, Nvidia Web page, 09 Nov 2017

<https://blogs.nvidia.com/blog/2016/07/29/whats-difference-artificial-intelligence-machine-learning-deep-learning-ai/>

Landscape overview

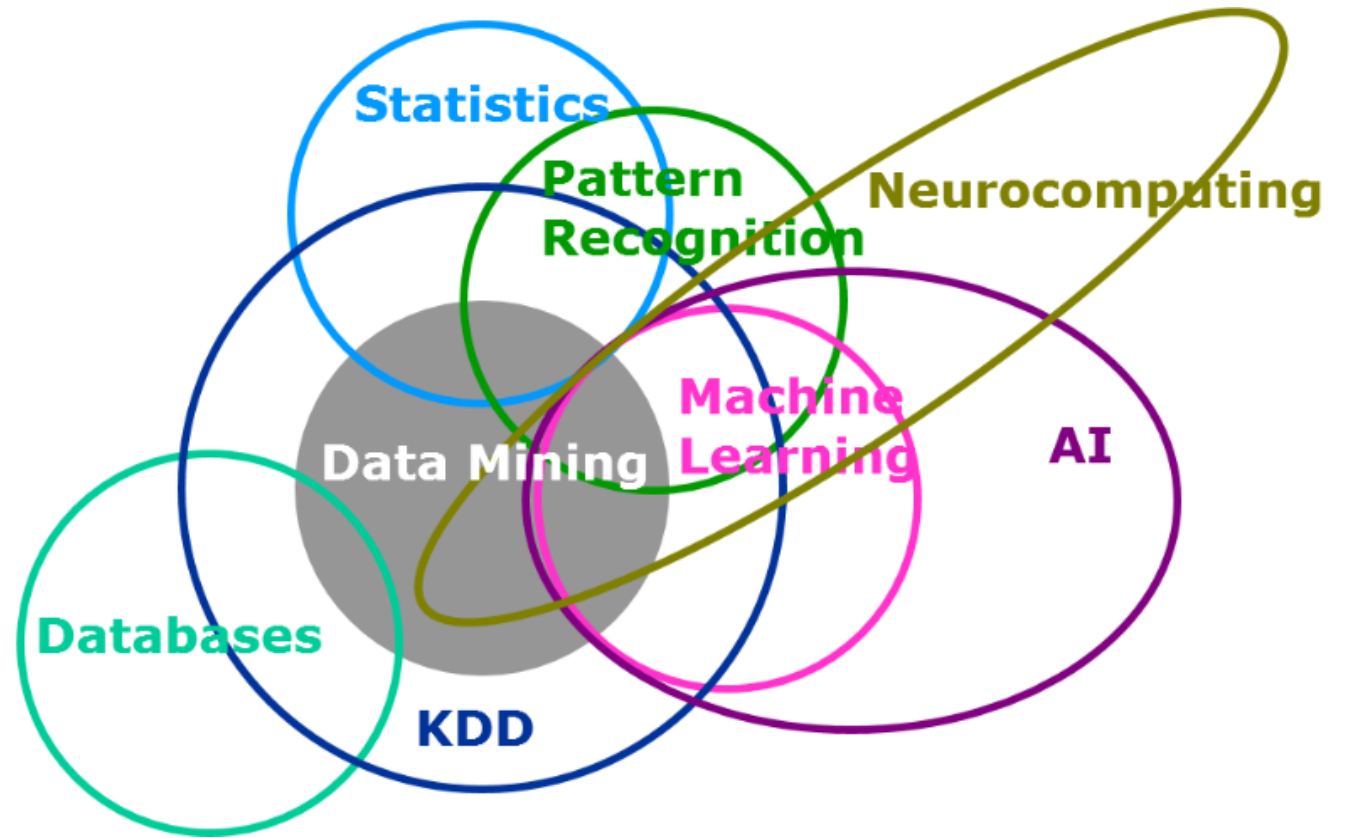
Landscape and focus² areas vary:

- Automation
- Statistics & Probability
- Data Mining
- Artificial Intelligence
- Machine Learning
 - Supervised Learning
 - Unsupervised Learning
 - Reinforcement Learning
- Deep Learning
- Other Areas
 - Deep Reinforcement Learning



Landscape overview

Landscape and focus³ areas of ML also could have overlaps with other scientific disciplines such as Data Mining, AI, Big Data and other.



Bayes Theorem

- $P(A|B) \implies$ Probability of A given that B

- Exercise on Bayes⁴

Out of 3000 emails received over a certain period, 2000 are spam and 1000 are not. The word “Rolex” appeared in 250 out of the 2000 which are spam and in 5 out of the 1000. So, if an email is received, and contains the word “Rolex”, what is the possibility that it is a spam?

Let S be the event that the message is spam, and E be the event that the message contains the word w . Under our assumption from before, we have that:

$$P(S|E) = \frac{P(E|S)}{P(E|S) + P(E|\bar{S})}$$

Bayes Theorem exercise

- Example – Solution:
Out of 3000 emails received over a certain period, 2000 are spam and 1000 are not. The word “Rolex” appeared in 250 out of the 2000 which are spam and in 5 out of the 1000.

So, if an email is received, and contains the word “Rolex”, what is the possibility that it is a spam?

$$P(S|E) = \frac{P(E|S)}{P(E|S) + P(E|\bar{S})}$$

$$P(S|E) = \frac{\frac{250}{2000}}{\frac{250}{2000} + \frac{5}{1000}} =$$

$$\frac{0.125}{0.125 + 0.005} \approx 0.962$$

Markov Chains

- Statistics & Probability - Markov Chains⁵


Markov Chains is a probabilistic process, that relies on the current state to predict the next state. For Markov chains to be effective the current state has to be dependent on the previous state in some way;

For instance, from experience we know that if it looks cloudy outside, the next state we expect is rain. We can also say that when the rain starts to subside into cloudiness, the next state will most likely be sunny.

MARKOV TABLE OF PROBABILITIES

STATE	NEXT STATE	PROBABILITY	%
CLOUDY	CLOUDY	0.1	10%
CLOUDY	RAIN	0.5	50%
CLOUDY	SUNNY	0.4	40%
RAIN	CLOUDY	0.3	30%
RAIN	RAIN	0.6	60%
RAIN	SUNNY	0.1	10%
SUNNY	CLOUDY	0.4	40%
SUNNY	RAIN	0.1	10%
SUNNY	SUNNY	0.5	50%

Markov State Diagram



Current State Vector

C	R	S
1	0	0

Figure 4

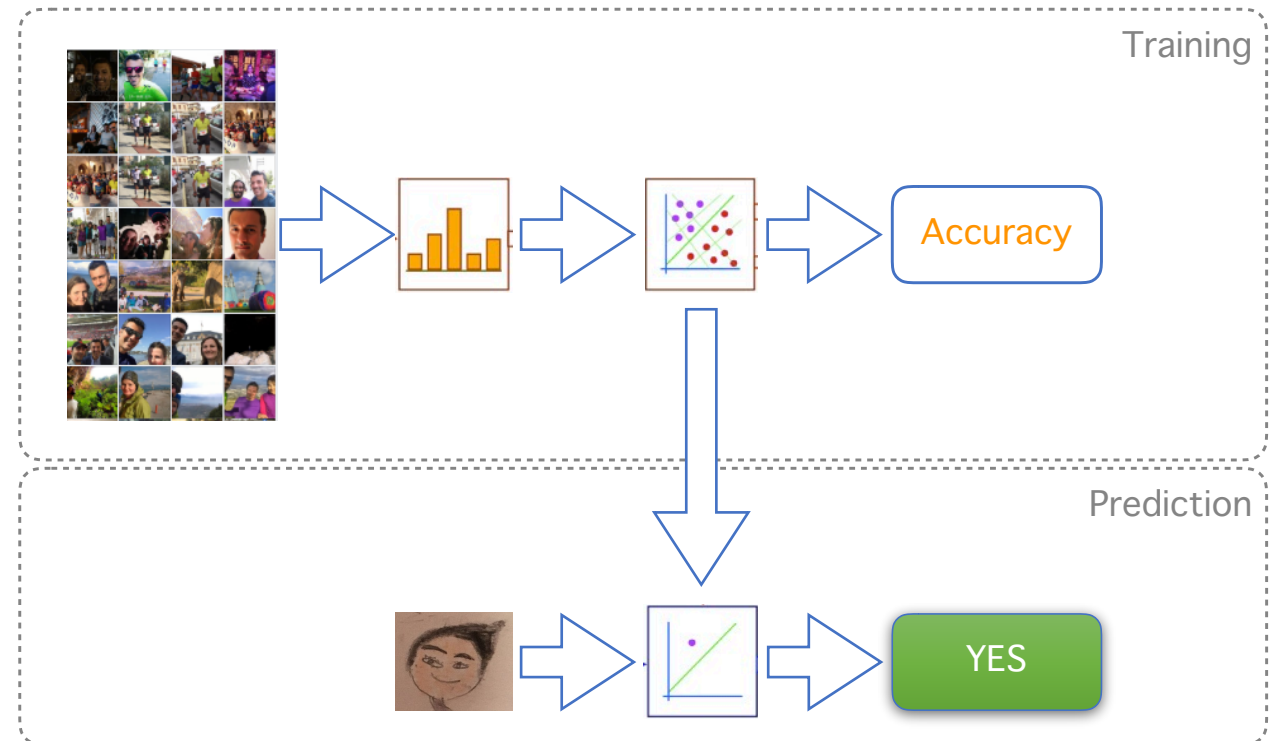
Statistical Classification⁶

- Examples of Classifiers

- Linear classifiers
- Support vector machines
- Quadratic classifiers
- Kernel estimation
- Boosting (meta-algorithm)
- Decision trees
- Neural networks
- Learning vector quantization

- Problems that classifiers help tackle:

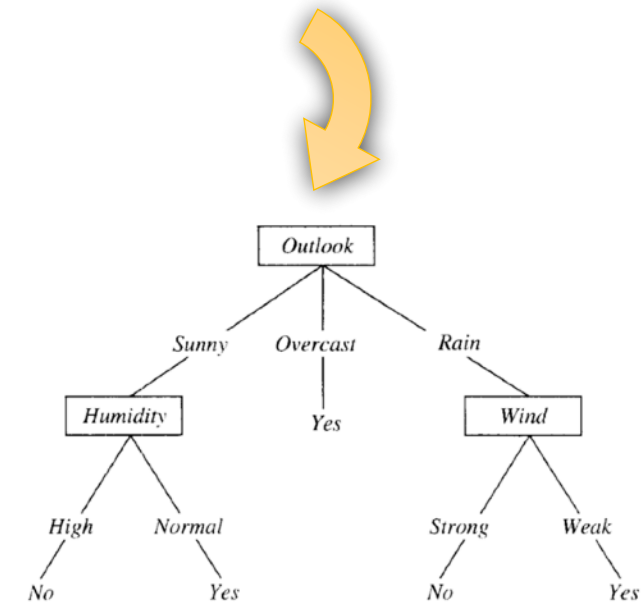
- Supervised Learning
- Clustering
- Dimensionality reduction
- Anomaly prediction
- Neural Nets
- RL/DRL



Supervised Learning - DTL

- Decision Tree Learning
DTL is method for approximating discrete valued target functions, in which the learned function is represented by a decision tree.
(Weka example will follow)
- Example^{7a} dataset converted via algorithm to Decision tree
- Methodology^{7b} of is whenever a feature is able to tell us more about our class, it is selected as a node

No.	1: outlook	2: temperature	3: humidity	4: windy	5: play
	Nominal	Numeric	Numeric	Nominal	Nominal
1	sunny	85.0	85.0	FALSE	no
2	sunny	80.0	90.0	TRUE	no
3	overcast	83.0	86.0	FALSE	yes
4	rainy	70.0	96.0	FALSE	yes
5	rainy	68.0	80.0	FALSE	yes
6	rainy	65.0	70.0	TRUE	no
7	overcast	64.0	65.0	TRUE	yes
8	sunny	72.0	95.0	FALSE	no
9	sunny	69.0	70.0	FALSE	yes
...	rainy	75.0	80.0	FALSE	yes
...	sunny	75.0	70.0	TRUE	yes
...	overcast	72.0	90.0	TRUE	yes
...	overcast	81.0	75.0	FALSE	yes
...	rainy	71.0	91.0	TRUE	no



7a Machine Learning, Mitchell, McGraw, 1997.

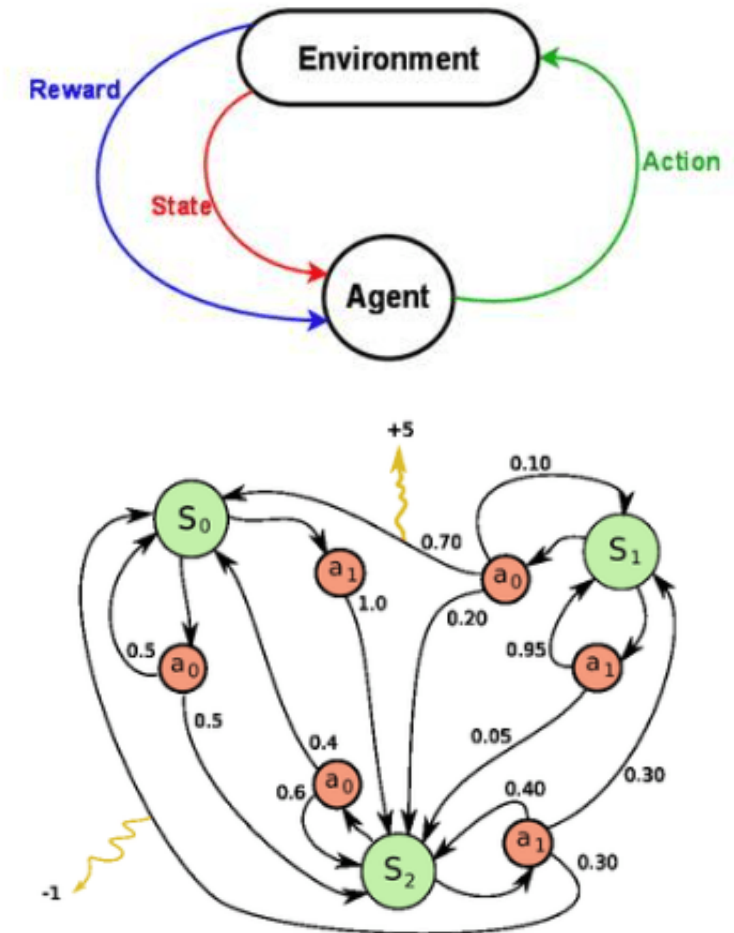
<http://www.cs.princeton.edu/courses/archive/spr07/cos424/papers/mitchell-dectrees.pdf>

7b Classification Methods, Padhye, 2017.

<http://www.d.umn.edu/~padhy005/Chapter5.html>


Reinforcement Learning (RL)

- Reinforcement Learning⁸ is learning what to do--how to map situations to actions--so as to maximise a numerical reward signal.
- Reinforcement learning is defined not by characterising learning methods, but by characterising a learning problem.



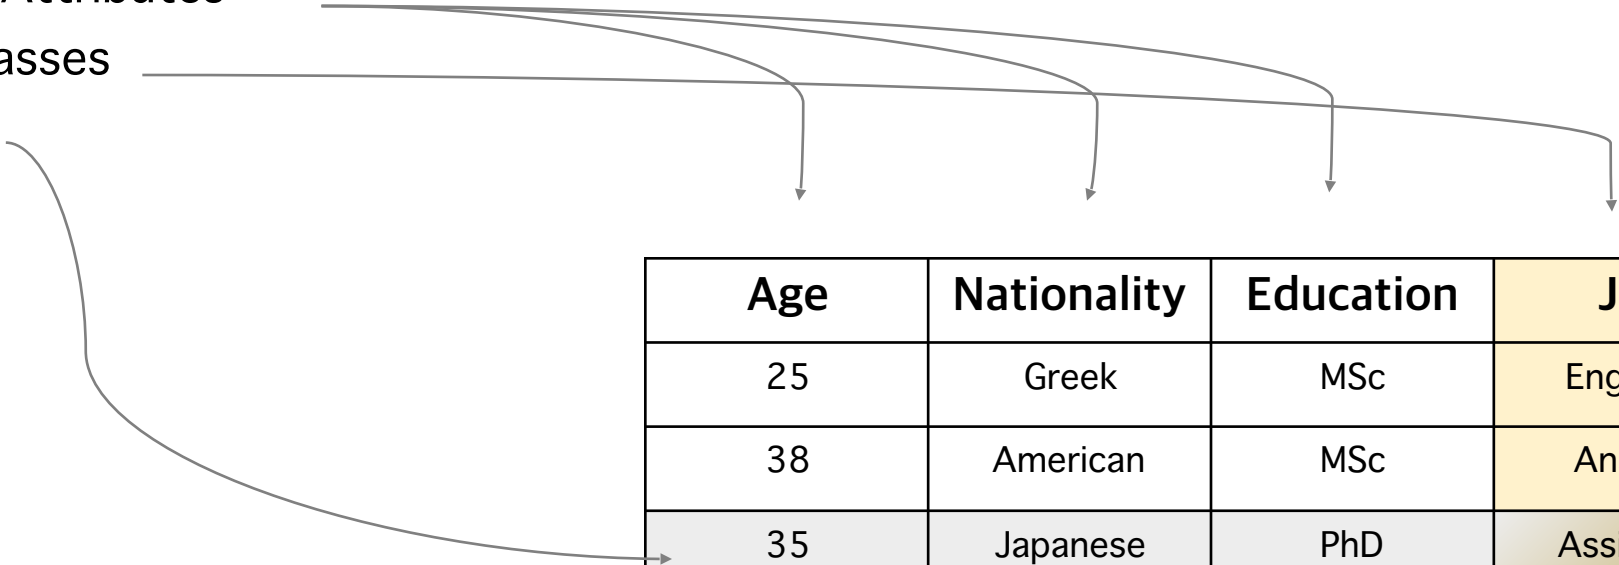
Workflow – The ML Pipeline⁹

- Start with the question or problem we want to solve
- Find proper data and sources, prepare data set (train/dev/test)
- Choose a model e.g. Decision Tree C4.5 Algorithm
- Train system & classify, Test
- Evaluate the system and fine-tune
- Predict / Forecast
- Apply to workflow
- Automate into workflow



Dataset semantics

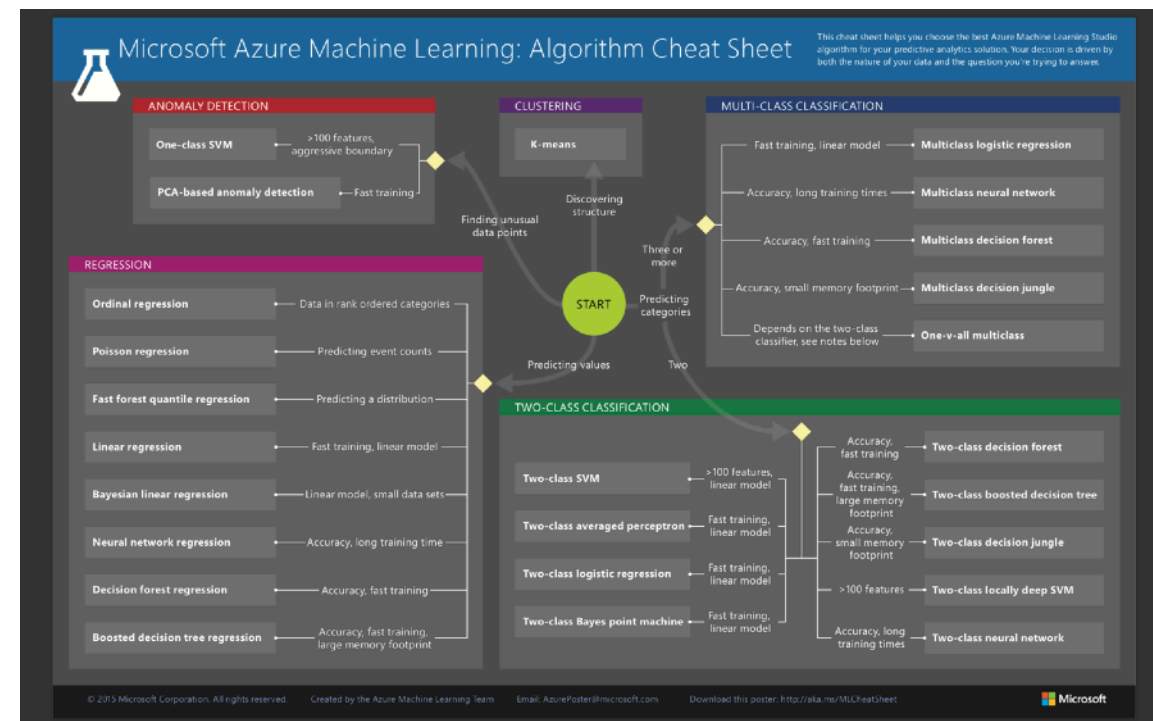
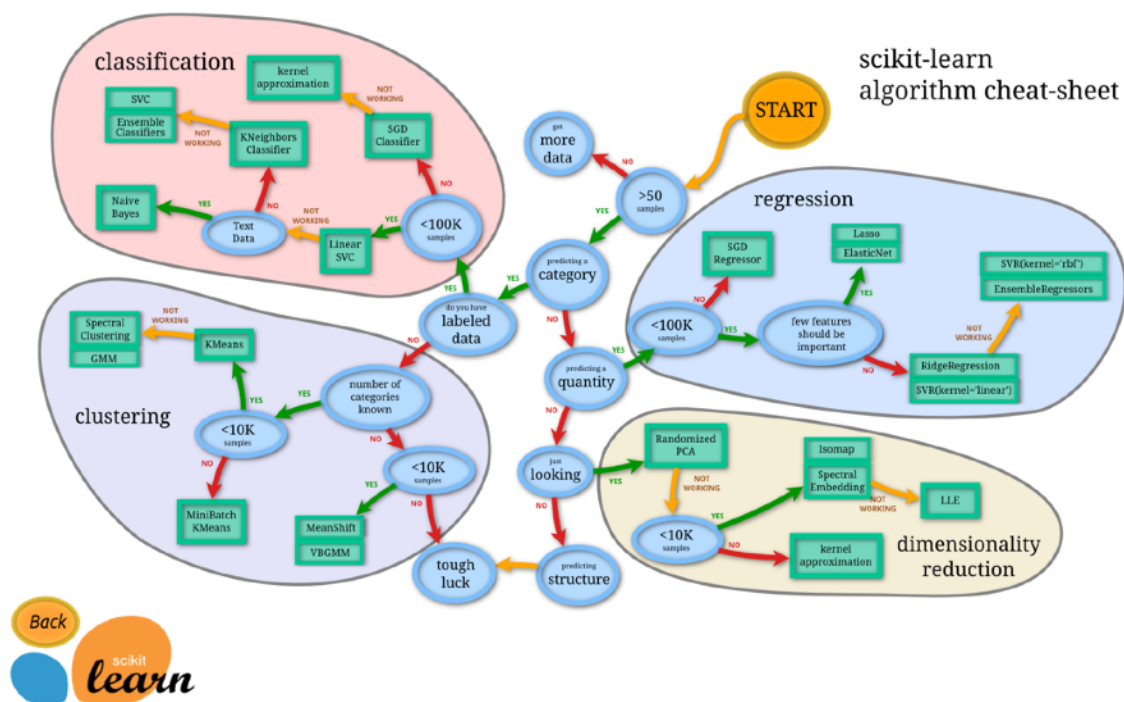
- Dataset
 - Features or Attributes
 - Labels or Classes
 - Instances



The diagram illustrates the mapping between the dataset components and the table structure. Arrows from 'Features or Attributes' point to the 'Age', 'Nationality', and 'Education' columns. An arrow from 'Labels or Classes' points to the 'Job' column. An arrow from 'Instances' points to the third row of the table, which represents a specific data instance.

Age	Nationality	Education	Job
25	Greek	MSc	Engineer
38	American	MSc	Analyst
35	Japanese	PhD	Assistant
...

Choose an algorithm^{10a, 10b}



10a Microsoft Azure Machine Learning: Algorithm Cheat Sheet, Microsoft website, 09 Nov 2017

<http://download.microsoft.com/download/A/6/1/A613E11E-8F9C-424A-B99D-65344785C288/microsoft-machine-learning-algorithm-cheat-sheet-v6.pdf>

10b Scikit-Learn Algorithm selection Procedure, Scikit-learn website, 23 Oct 2017

http://scikit-learn.org/stable/tutorial/machine_learning_map/index.html

Tools – Landscape¹¹

Category	Sub-category	Description	Examples
Packages of ML Implementations	Statistical Software Packages	Software toolkits with a large set of implementations of ML algorithms, typically with visualization support	SAS, R, Matlab, SPSS
	Data Mining Toolkits	Software toolkits with a relatively limited set of ML algorithms, typically over a data platform, possibly with incremental maintenance	Weka , AzureML, ODM, MADlib, Mahout, Hazy-Classify
	Developability-oriented Frameworks	Software frameworks and systems that aim to improve developability, typically from academic research	GraphLab, Bismarck, MLBase
	SRL Frameworks	Implementations of statistical relational learning (SRL)	DeepDive
	Deep Learning Systems	Implementations of deep neural networks	Google Brain, Microsoft Adam
	Bayesian Inference Systems	Systems providing scalable inference for Bayesian ML models	SimSQL, Elementary, Tuffy
Linear Algebra-based Systems	Statistical Software Packages	Systems offering an interactive statistical programming environment	SAS, R, Matlab
	R-based Analytics Systems	Systems that provide R or an R-like language for analytics, typically over a data platform, possibly with incremental maintenance	RIOT, ORE, SystemML, LINVIEW
Model Management Systems		Systems that provide querying, versioning, and deployment support	SAS, LongView, Velox
Systems for Feature Engineering		Systems that provide abstractions to make feature engineering easier	Columbus , DeepDive
Systems for Algorithm Selection		Systems that provide abstractions to make algorithm selection easier	MLBase, AzureML
Systems for Parameter Tuning		Systems that provide abstractions to make parameter tuning easier	SAS, R, MLBase, AzureML

¹¹ A Survey of the Existing Landscape of ML Systems, Kumar; McCann; Naughton; Patel, 27 Nov 2015

Deep Learning Toolkits comparison¹²

Toolkit	GPU Support	Other
Caffe	Yes	JSON-like text file to describe the network architecture
Deeplearning4j	Yes	Java on Scala API
Tensorflow	Yes	Google backing, high adoption - Python
Theano		Python
Keras		Python - uses Theano or Tensorflow as backend
MXNet	Yes	C++
Lasagne		Python - uses Then
CNTK		VS for ML - developed by Microsoft
DIGITS		Nvidia - web based tool
Torch		Written in C
PyTorch	Yes	Python frontend
Pylearn2		Python
Chainer		

¹² Toolkits and Libraries for Deep Learning, Erickson; Korfiatis; Akkus; J Digit Imaging 30: 400., 2017

<https://doi.org/10.1007/s10278-017-9965-6>

Before-Selecting-a-Tool Checklist

- Things to consider for a toolkit/tool/ecosystem
 - Environment ease of use
 - Dev & Exec speed
 - Training Speed
 - GPU Support
 - Community support & contributors
 - License contamination
 - Language to be used



ANACONDA

PYTHON & R OPEN SOURCE ANALYTICS

NumPy	SciPy	Pandas	Scikit-learn	Jupyter/IPython	
Numba	Matplotlib	Spyder	TensorFlow	Cython	Bokeh
Scikit-image	NLTK	Dask	Caffe	dplyr	shiny
ggplot2	tidyr	caret	PySpark	& 1000+ packages	
 CONDA					



R Studio



elastic




Practice (fun part :-)

- Weka - Statistics - Pipeline - DTL
- R Studio - Statistics
- Orange Data Mining - DTL
- Anaconda Python - Statistics - basic ML
- Elasticsearch (ELK) - Visual Analytics

on:

- bitcoin
- cars
- flights
- milano_cells
- predictive_maintenance



Lectures & Data Sources Page
<https://github.com/sgez/MLAI>

Thank you!

“In theory there is no difference between theory and practice. In practice there is.”

Yoggi Berra